

Answer: In this plot we have the partially true binary classifier labels y for training a logistic regression classifier. Once again the plot is on a withheld test dataset, and the points' classes are of the true labels. We can see the decision boundary is shifted so that both positive (blue cross) and negative (green dot) labelled datapoints lie on the negative half - this corresponds to the fact that the y labels are all true positives, true negatives, false negatives, but no false positives, i.e. we're biased towards classifying more points as negative.

